

A desperate guide to data analysis workflow improvements



Peter Regner

“Reproducible Data Analysis”
Data Science @ BOKU initiative

2024-02-29
BOKU Wien



THIS IS YOUR MACHINE LEARNING SYSTEM?


YUP! YOU POUR THE DATA INTO THIS BIG PILE OF LINEAR ALGEBRA, THEN COLLECT THE ANSWERS ON THE OTHER SIDE.

WHAT IF THE ANSWERS ARE WRONG?

JUST STIR THE PILE UNTIL THEY START LOOKING RIGHT.



ME TRYING TO EXPLAIN




MY DATA PROJECT

Our goal: reproducibility!

...but in practice:

nobody re-runs it
or it is broken



irreproducibility
by hardware
demands




irreproducibility
by obscurity

traceability



comprehensibility



some users
might be
interested
mostly in
results

Make things more usable - for yourself and others!

Reproducibility in practice: How easily can somebody use the thing later?

Think about your target audience to set priorities!

Do they want to...

- re-run all your scripts?
- run the code as piece of software?
- use the code and adapt it?
- use the computation results only?
- understand what and how we did it?



7 Steps towards more use-able data analysis projects

Gamification: count every ✓, if you follow the guideline already!



(1) Publish your result data sets ✓

Wicherts, J. M., Bakker, M., & Molenaar, D. (2011). Willingness to Share Research Data Is Related to the Strength of the Evidence and the Quality of Reporting of Statistical Results. In R. E. Tractenberg (Ed.), PLoS ONE (Vol. 6, Issue 11, p. e26828). Public Library of Science (PLoS). <https://doi.org/10.1371/journal.pone.0026828>

Willingness to Share
Research Data

is related to

Strength of
the Evidence

(1) Publish your result data sets: in a fair way!


F
indable




A
ccessible



I
nteroperable



R
eusable



<https://www.go-fair.org/fair-principles/>
<https://doi.org/10.1038/sdata.2016.18>

by [SanjyaPundir](#) under CC-BY-SA 4.0

(1) Publish your result data sets: how?

- ✓ Use an open license if possible! (Allows others to publish your results as input data!)
- ✓ Label and document the fields of your result data set properly
- ✓ A CSV or NetCDF file on Zenodo is already pretty good
- ✗ tables in the PDF of the supplementary material
- ✗ non-open or non-standard file formats like *.gdx (GAMS)
- ✗ units for quantities not documented

More data providers: Dataverse, FigShare, AUSSDA, Dryad, Mendeley Data, DataHub, DANS, EUDat, ...

(2) Publish your code ✓

- use an open license, e.g. MIT license or CC-BY ✓ (unless you need GPL libraries)
- For maintained projects: Use GIT / Github ✓
- Include a README file ✓
 - What does it do? (summary & link to article) ✓
 - Requirements: hardware & software ✓
 - How to run / use the thing ✓

(3) Use reproducible virtual environments ✓

Virtual environments (e.g. Conda) allow you to:

- Multiple versions of a library on one machine
- Pretty portable (Windows, Linux, Mac OS)
- Easier to use than Docker

...and to export all used libraries & precise versions!

Alternatives to conda: renv, pipenv, ...

Export all used dependencies using conda: ✓

```
conda env export --no-builds | grep -v "^prefix: " > env.yml
```

Use micromamba - conda is too slow! ✓

<https://mamba.readthedocs.io/en/latest/installation/micromamba-installation.html>

env.yml:

```
name: my_project
channels:
  - conda-forge
  - defaults
dependencies:
  - ca-certificates=2023.11.17
  - liblapack=3.9.0
  - libzlib=1.2.13
  - ncurses=6.4
  - numpy=1.26.3
  - openssl=3.2.1
  - pip=23.3.2
  - python=3.10.13
  [...]
  - wheel=0.42.0
  - xz=5.2.6
  - zstd=1.5.5
  - pip:
    - cplex==22.1.1.0
```



(4) Automate everything

Manual steps are...

- error prone
- time consuming
- not documented

Every command or click should be stored in a way, such that it can be executed again!

Be realistic! You will do things over and over again...



But you can automate all steps!

NOTHING GOOD

Please don't
take this too
literally!



HAPPENS IN EXCEL...!

(4) Automate everything

Goal: a single executable script ✓

- download of input data
- preprocessing
- computation
- write results to files
- create figures

article

L^AT_EX

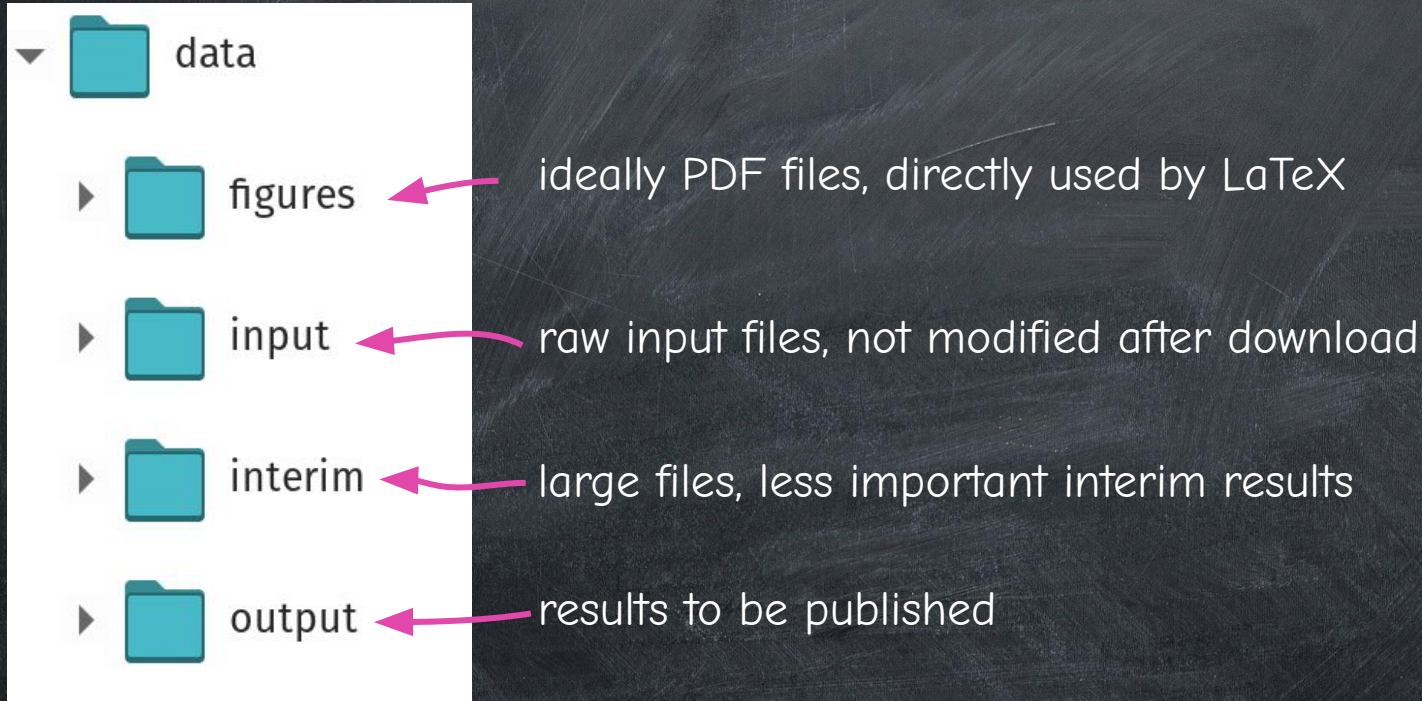
or



how to &
links in
bonus slides

The screenshot shows a GitHub repository interface. A file named 'run.sh' is highlighted with a pink oval. Below the file list, there are links for 'README' and 'MIT license'. A status bar shows 'license MIT', 'Tests passing', and 'DOI 10.1088/2515-7620/ace0b9'. The main content area displays the title 'Explaining the decline of U density'.

(5) Use a good folder structure ✓



(5) Use a good folder structure

Follow common conventions!

- no spaces in file names ✓
- file names in lower case or ALL_CAPS ✓
- avoid CamelCase for files and folders ✓



<https://github.com/cookiecutter/cookiecutter>

Literature:

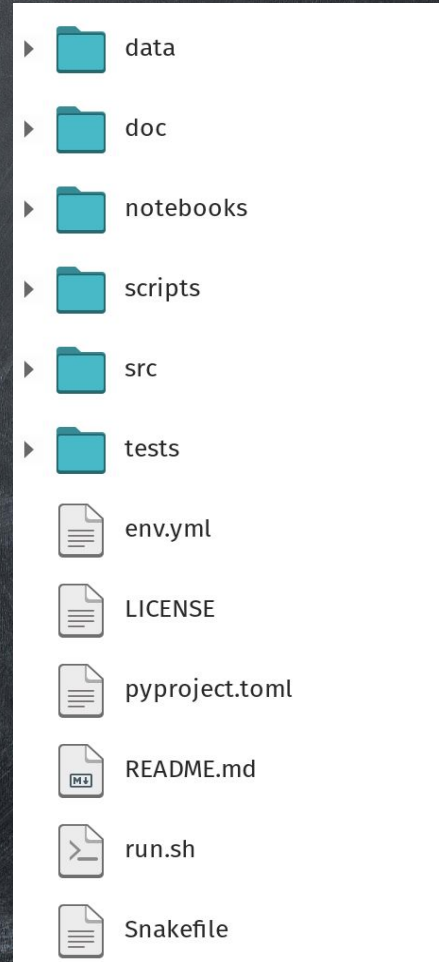
[1] G. Wilson, J. Bryan, K. Cranston, J. Kitzes, L. Nederbragt, and T. K. Teal, "Good enough practices in scientific computing," PLOS Computational Biology, vol. 13, no. 6, p. e1005510, Jun. 2017, doi: 10.1371/journal.pcbi.1005510.

[2] W. S. Noble, "A Quick Guide to Organizing Computational Biology Projects," PLOS Computational Biology, vol. 5, no. 7, p. e1000424, Jul. 2009, doi: 10.1371/journal.pcbi.1000424.


[3] "Cookiecutter Data Science." <https://drivendata.github.io/cookiecutter-data-science/>

your project


auto-create
from templates




Organize your computation pipeline



inputs



outputs

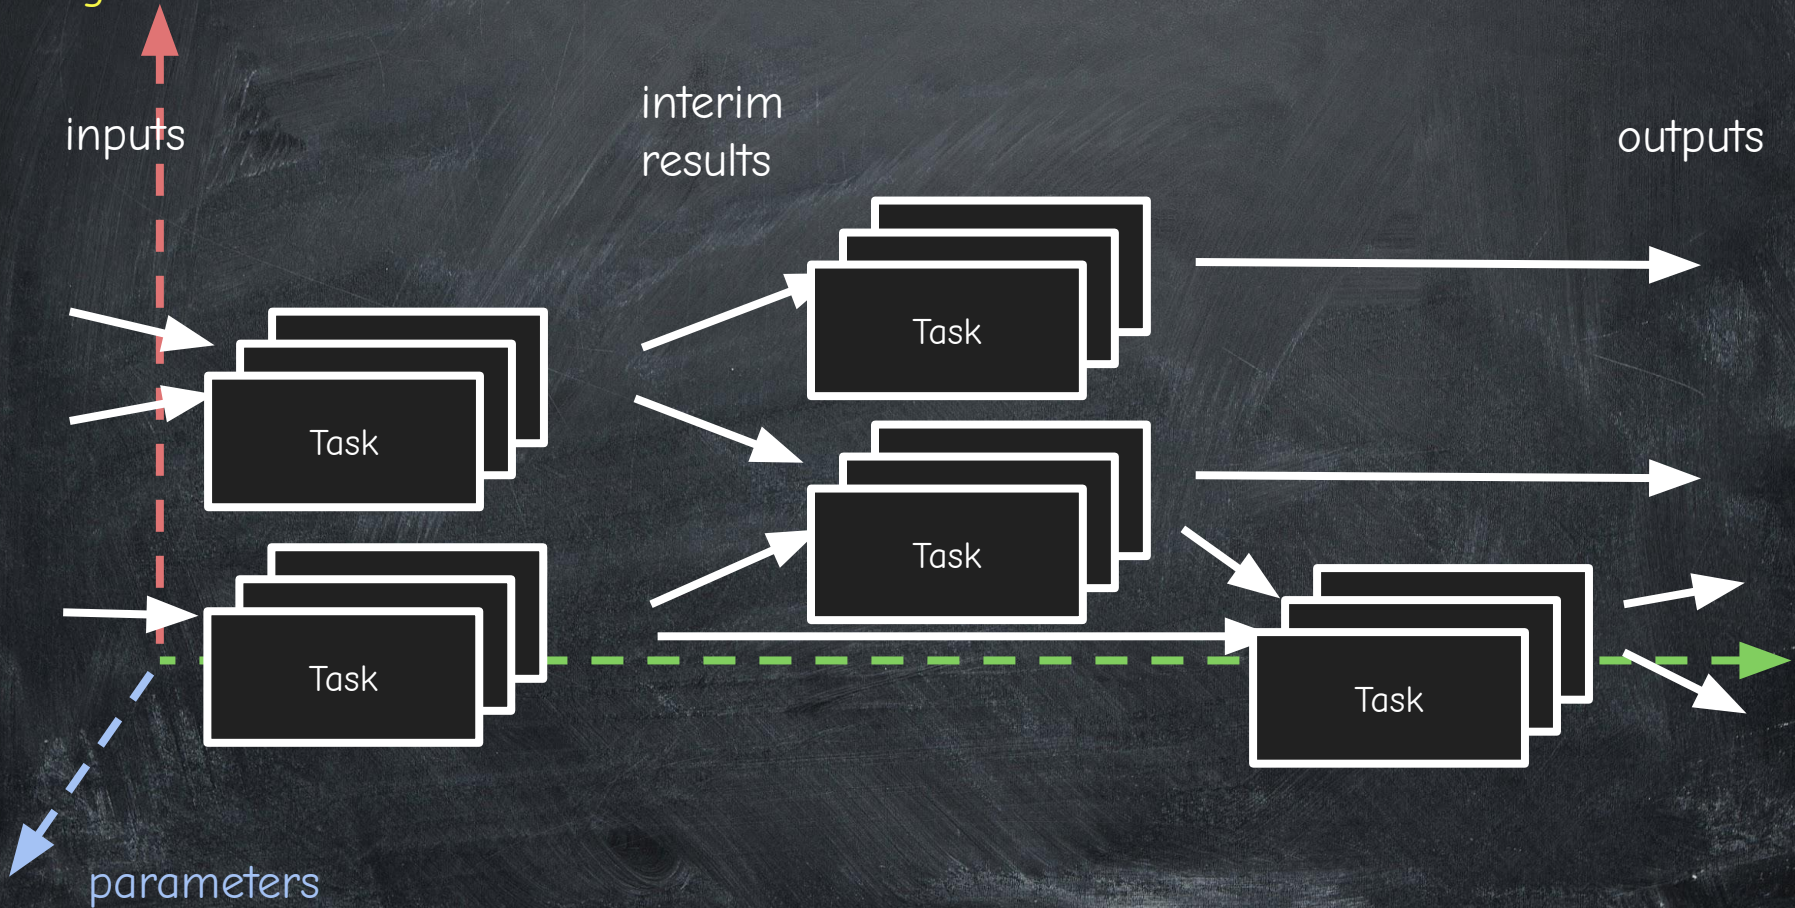


config

inputs

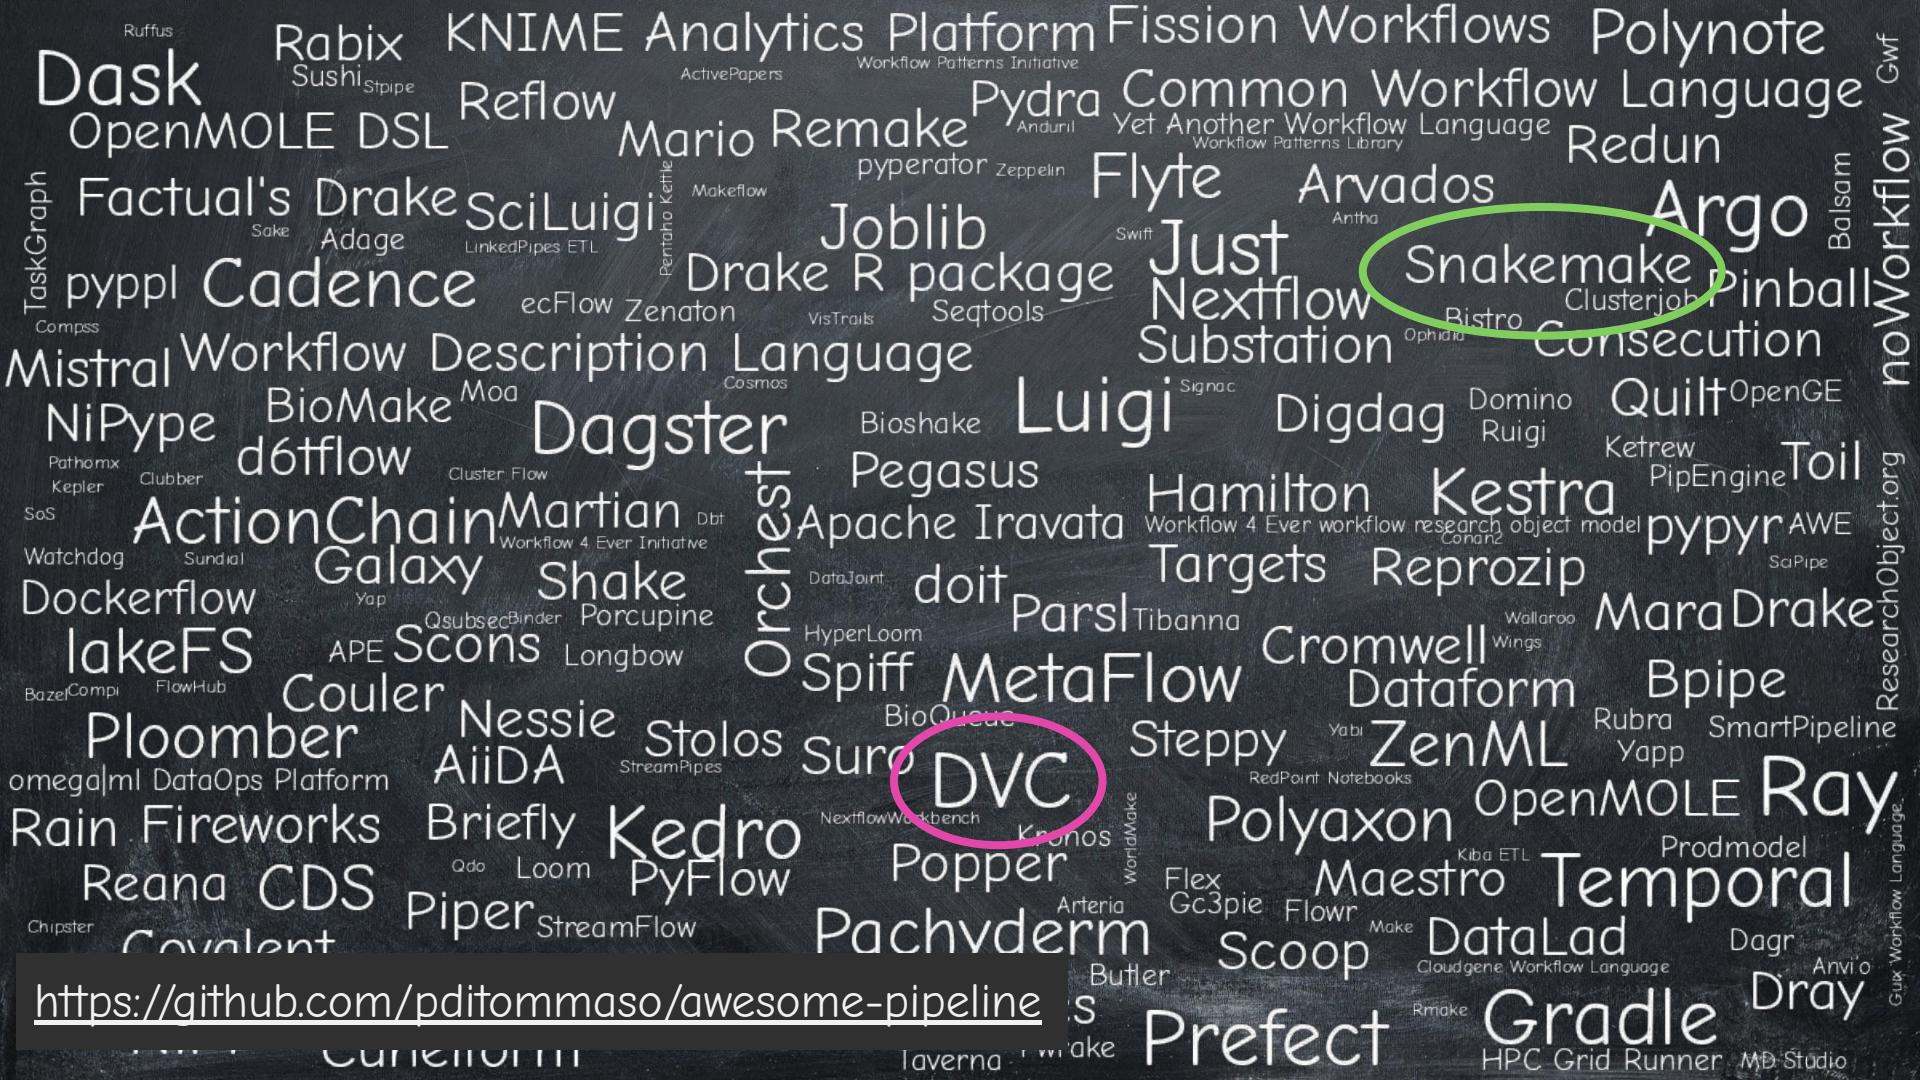
interim
results

outputs



(6) Use a pipeline tool ✓

- Caching / incremental builds
 - skip tasks if nothing changed (inputs, code, config, parameters)
- Parametrizing tasks
 - specify (multi dimensional) parameter spaces via ranges or lists
- Run tasks in parallel
 - on different CPUs or on different nodes on a cluster (e.g. VSC)



<https://github.com/pditommaso/awesome-pipeline>

(7) Keep old versions of your data ✓

Re-running computations might be difficult or take long time... Keep old versions!

It would be nice to...

- know how artefacts (figures, output files, ...) were generated
→ record Git hash / code, all inputs, parameters, conda environment, config, ...
- delete large files later when disk space is an issue
- compare between different versions of data

Git alone is not enough to fulfill these requirements!

Note that Git offers more features:

- Sync to other machines / collaboration with others
- Publish the repository - but data is not always publishable due to licence issues

Tools for large data: DVC, GIT-LFS, git-annex, DataLad, ...

My poor-man data versioning

run `.sh` script:

- runs the computation
- creates a Git tag
- copies outputs to an archive folder

Open issues:

- How to sync data between machines? Use Unison?
- How to collaborate with others using the same data?
- How to sync figures to Overleaf without Git?

<https://github.com/inwe-boku/windpower-decomposition-usa/blob/main/run.sh>



Cookies for all of you!




Everything we have seen today in a cookiecutter template:

- folder structure
- code snippet for including results in LaTeX
- Script for building a zip file for arXiv.org
- run.sh script
- Snakemake file
- README.md, .gitignore, LICENSE, ...
- ...

→ <https://github.com/inwe-boku/cookiecutter-data-research>

Reproduce all the things!

Download slides:



<https://bit.ly/desperate-guide>

peter.regner@boku.ac.at

<https://github.com/inwe-boku/>

<https://refuel.world/>



We gratefully acknowledge support from the European Research Council ("reFUEL" ERC-2017-STG 758149).

reFUEL

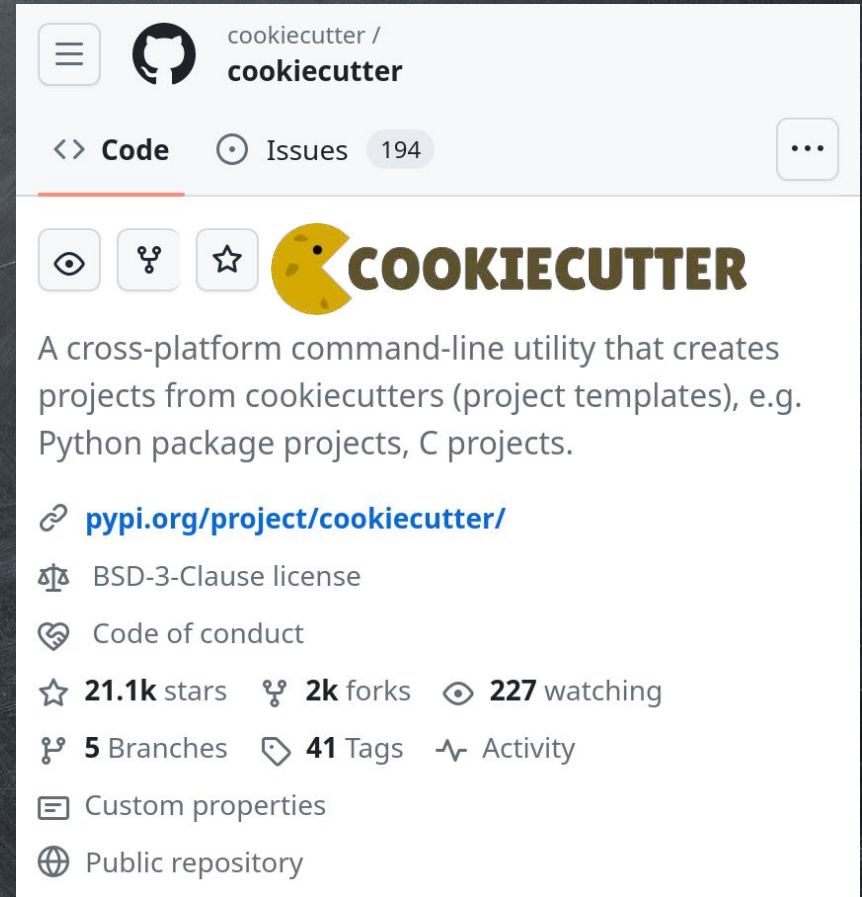


Automate project creation ✓

Create a new folder from a template:

- folder structure
- LICENSE file
- .gitignore file
- code snippets
- ...

<https://github.com/cookiecutter/cookiecutter>



The screenshot shows the GitHub repository page for `cookiecutter / cookiecutter`. The repository is a public repository with 194 issues. It features a yellow Pac-Man logo and the text **COOKIECUTTER**. The description states: "A cross-platform command-line utility that creates projects from cookiecutters (project templates), e.g. Python package projects, C projects." The repository is licensed under BSD-3-Clause and has a code of conduct. It has 21.1k stars, 2k forks, and 227 watchers. There are 5 branches and 41 tags. The repository is linked to pypi.org/project/cookiecutter/.

Include figures in LaTeX documents ✓

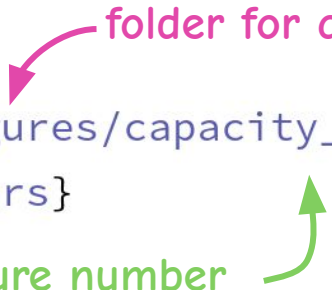
MS Word: use linked images

Latex:

```
\begin{figure}
  \centering
  \includegraphics{data/figures/capacity_factors.pdf}
  \label{fig:capacity_factors}
\end{figure}
```

folder for all figures

no figure number



Overleaf: use GIT, Github or dropbox to sync figures - premium only... :-/

https://www.overleaf.com/learn/how-to/Dropbox_Synchronization

https://www.overleaf.com/learn/how-to/Git_Integration_and_GitHub_Synchronization

Include results in LaTeX documents ✓

Turns out to be surprisingly difficult! 😱

```
result_values = {}  
  
meaning_of_life = 42  
result_values['meaning_of_life'] = f'{meaning_of_life:d}'  
  
gravity_ms2 = 9.80665  
result_values['gravity'] = f'{gravity_ms2:.2f}', 'm/s^2'  
  
write_result_values(result_values)
```

script writes results to a file with custom LaTeX commands

```
\newcommand{\meaning_of_life}{42}  
\newcommand{\gravity}{\qty{9.81}{m/s^2}}
```

auto-generated LaTeX commands

More details and other Solutions like Knitr:
<https://tex.stackexchange.com/a/711627/8964>

Do more good things

- Write good code
 - Don't do no magic numbers ✓
 - No absolute file paths ✓
 - No unnecessary Abbrev. ✓
 - Follow code conventions (e.g. pep8) ✓
- Write unit / functional tests ✓

Automatically check if things are behaving as expected!
- Use a code linter (e.g. as GIT hook) ✓

Tools to automatically check code for errors and style violations, for Python: flake8, black (auto-formatter),
- Use Continuous integration ✓

Run tests, code linter or computations automatically, e.g. via Github Actions
- Do code review ✓

A great way to improve quality and spread knowledge in your team!

Tracking additional computation information

@task

```
def my_fancy_func(some_param, inputs, outputs):  
    # do something here  
    ...
```

The @task decorator creates a file with metaparameters for each output:

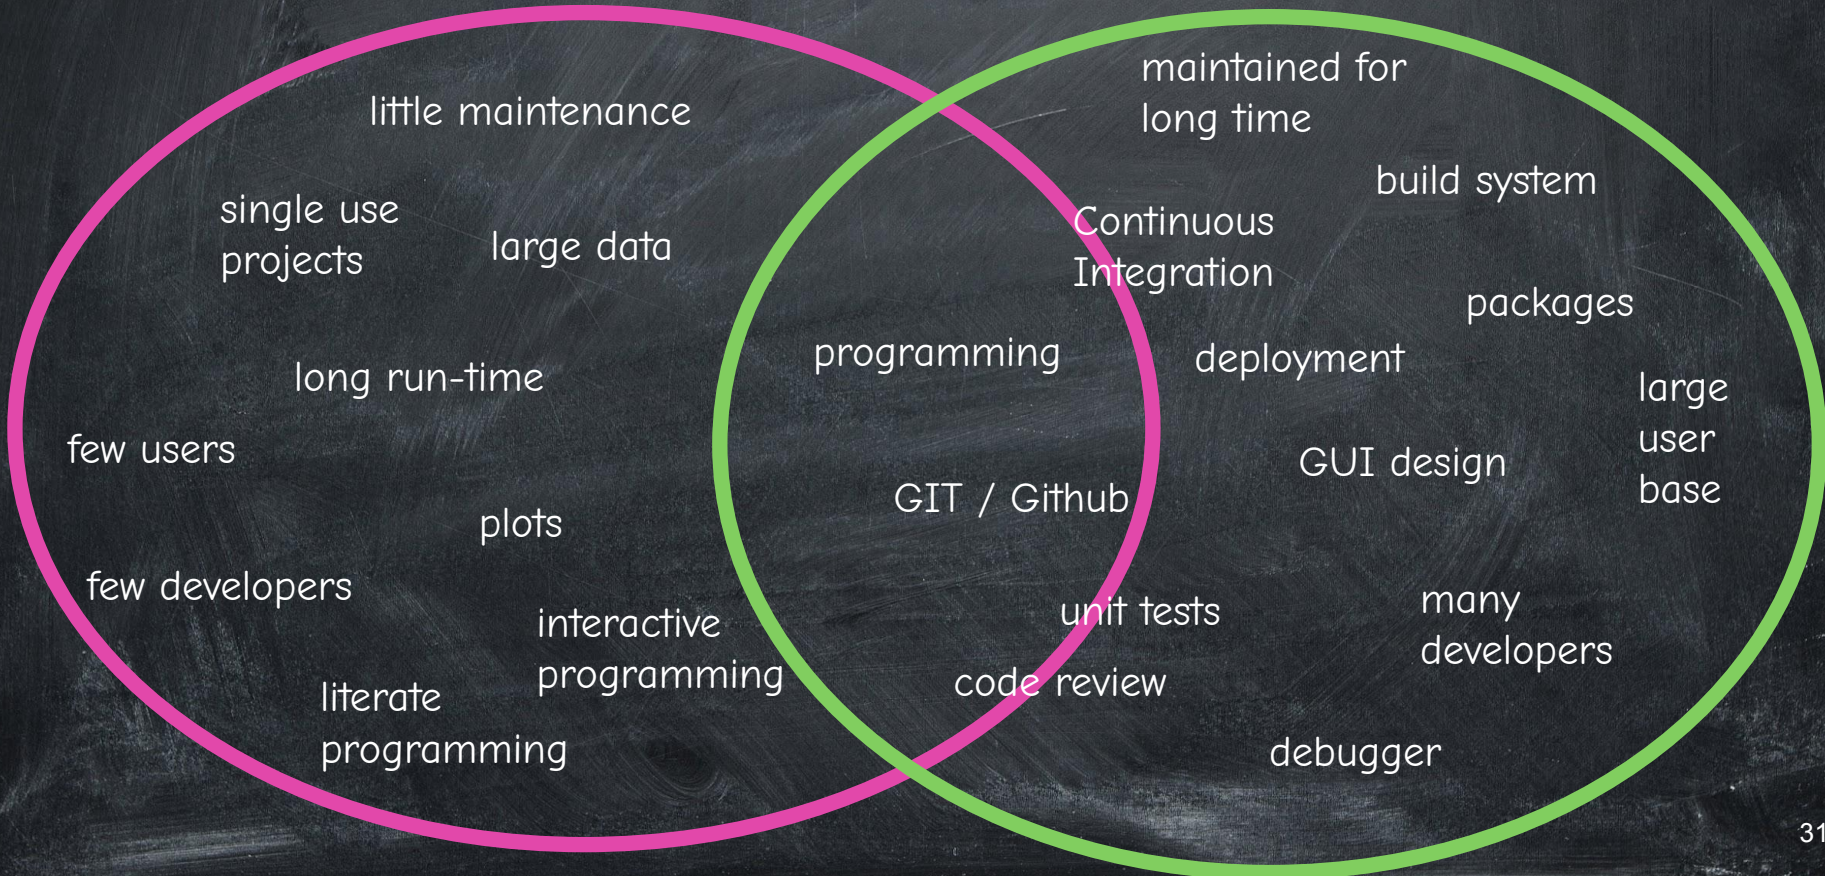
```
function: concat_solution_chunks  
git_commit: 59f9737-dirty  
hostname: nora  
input_files:  
- data/interim/network_solution/network_solution_740_560.nc  
  [...]  
- data/interim/network_solution/network_solution_795_615.nc  
input_params: {}  
output_files:  
- data/output/network_solution/network_solution.nc  
runtime: 7.013682842254639  
start_time: '2024-01-10T12:57:21.699562+01:00'
```

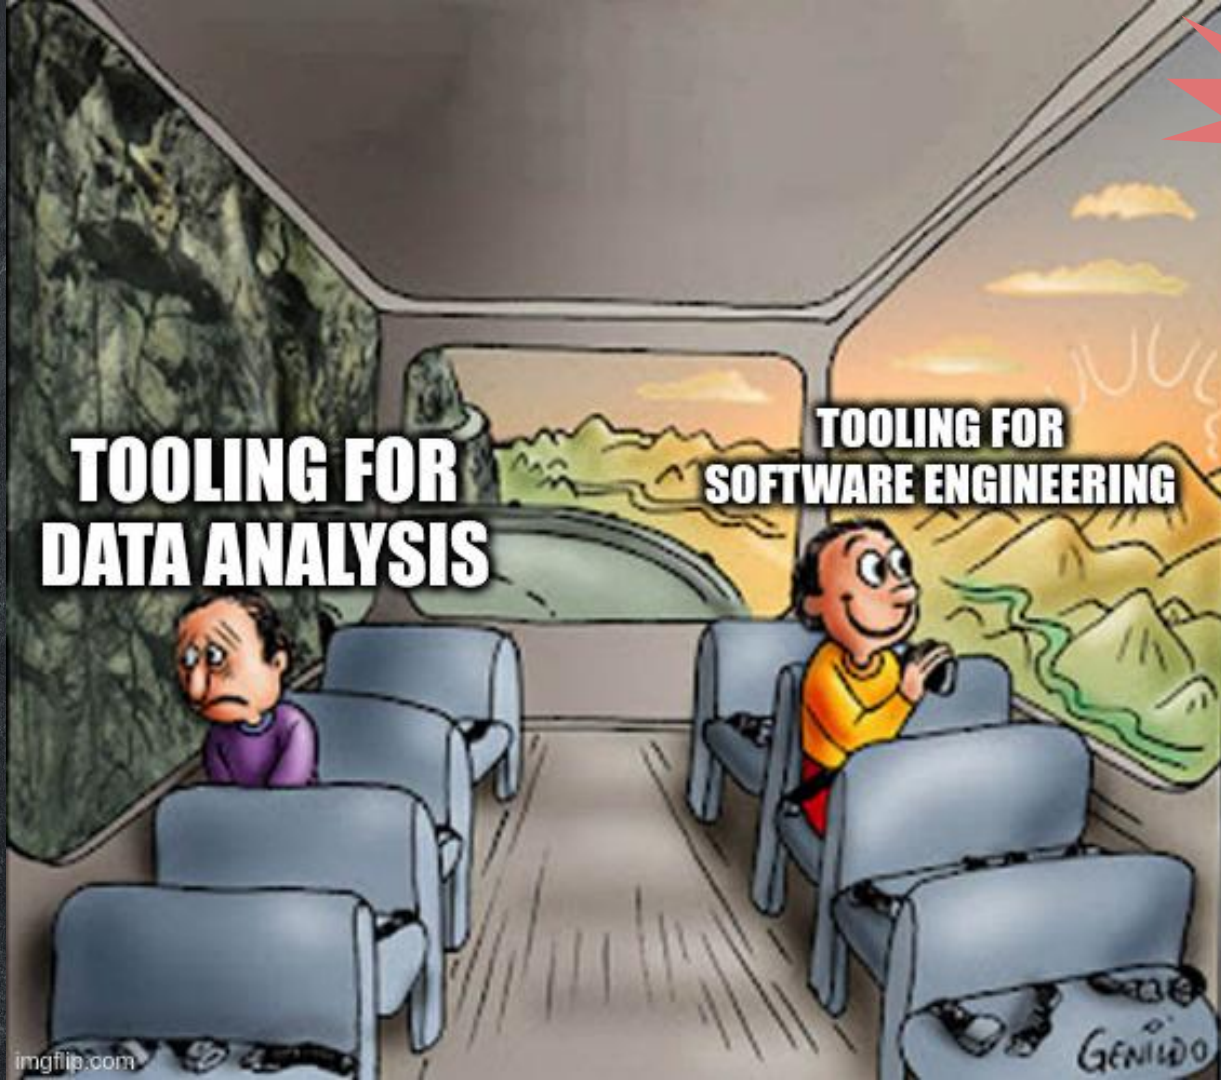
Link to source:

https://github.com/inwe-boku/cookiecutter-data-research/blob/main/%7B%7B%20cookiecutter_epo_name%20%7D%7D/src/task.py

Data Analysis

Software Development





**TOOLING FOR
DATA ANALYSIS**

**TOOLING FOR
SOFTWARE ENGINEERING**

opinionated!

Reproducibility is a hot topic...



Toni Whited 🇺🇦 🤖
@toniwhited

Texting with a colleague at another uni who is mired in the AEA replication process. Wasting > month getting all in order so that some RA can just push one button and get all the data merges right so that the results are perfect to the 3rd decimal place.

1/5

6:05 PM · Feb 17, 2024 · **429K** Views



It should not be the only priority, but I'd be a bit skeptical about code quality if it takes more than a month to get this done.

<https://twitter.com/toniwhited/status/1758900448596291953>